

Aims & Objectives

University students routinely face content overload such lecture decks, PDFs and readings pile up faster than they can be broken down into actionable study material. Time that could be spent on retrieval practice is instead swallowed by copying slides into flash-cards, hunting for definitions and devising self-tests. The result is low retention, last-minute cramming, and stress, resulting with outcomes at odds with research on spaced repetition and cognitive-load theory.

SkillSprint will tackle this problem by acting as a one-click personal tutor, web tool that accepts any slide deck or PDF and within seconds outputs a sequenced micro-lesson comprising flash-cards, a concise reading summary and Bloom-aligned quiz questions optimized for sixty-second study intervals. The primary aim is to empower time-pressed learners to convert raw material into evidence-based micro-learning units that boost recall while cutting down on preparation effort.

To make that aim concrete, the project sets five main objectives:

1. Cut preparation time by at least 50 % (import → first quiz) compared with manual flash-card creation, as measured in a formative baseline study by Week 6.
2. Deliver a user experience that achieves a System Usability Scale (SUS) score of 80 or higher (the threshold for “good/excellent”) by Week 8.
3. Reach an average quiz-item accuracy of 80 % or better when questions are judged by subject-matter experts against the source material, by project submission.
4. Produce a 20 % or greater improvement in delayed recall (immediate vs. 48-hour post-test) for at least 70 % of evaluation participants ($n \geq 20$) by project submission.
5. Keep median generation latency at or below five seconds for a 20-slide deck on standard consumer hardware, verified by Week 4.

These quantifiable targets double as pass/fail thresholds for later User-Acceptance Tests and will drive backlog prioritisation and sprint planning in the Agile schedule that follows.

Planning and Agile Method

To translate the five SMART objectives into executable work, the project is organized as four two-week sprints plus a one-week buffer. Each sprint must end with :

- (a) A usable slice that a student tester can try
- (b) One concrete metric captured (latency, SUS, accuracy, or recall).
- (c) A single Trello board holds every story; columns are Backlog → Sprint → In Dev → Review/Test → Done. Velocity is measured after Sprint 1 and used to adjust the story load of later sprints.

Weeks	Sprint Goal	Objective	Key Metric Checkpoint
1-2	Foundation & Import Engine	PDF/slide parser that extracts headings & paragraphs into JSON. Create diagrams, send user-survey	Documenting Risk logs, to identify flaws and challenges early on,

Weeks	Sprint Goal	Objective	Key Metric Checkpoint
3-4	Quiz-Generator Core	Prompt-chain that turns a JSON chunk into at least one validated MCQ set	≤ 5 s median latency (to generate the results)
5-6	Usability & Micro-Lesson Flow	Usable prototype integrated to live API. First round of SUS sent to users.	$\geq 50\%$ time saved
7-8	Refinement & Evaluation	Improved UI, onboarding. Second round of SUS sent to users. Audit of quiz accuracy.	SUS ≥ 80 % and Accuracy of model $\geq 80\%$
9	Contingency Plans	Polish documentation, Create Final User Acceptance Testing Matrix	-

Backlog Structure & Traceability

Epics (in scrums) will mirror the rubric sections (e.g. Import, Generation, UX, evaluations, Ops). Stories are written as user stories with acceptance criteria that reference the objectives or User Acceptance Testing items. An example can be seen like the following:

Story 3.2 – As a learner I want to receive ten flash-cards generated from my slides

Acceptance: Given a 20-slide deck, when I click “Generate” then ≤ 5 s later a flash-card list appears with ≥ 80 % content relevance.

The UAT matrix (§ 3.4) is auto-generated from these acceptance tests via a small Python script so that specification and backlog stay in sync.

Risk-managed capacity

he Trello board has a red “Risk” label. Any story tagged risk must enter the sprint no later than week 4. Current top items are:

Risk Label	Mitigation
LLM API cost	Cache quiz results, keep an open-source fallback branch.
Complex PDF Layouts Break Parser	Maintain a 30-file edge-case corpus, buffer week reserved for fixes.
SUS < 80 after round 1	Schedule a rapid heuristic review and design workshop mid-Sprint 5; prioritise usability fixes over new features.

Risk Label	Mitigation
Data-privacy concerns over uploaded files.	Client-side parsing, include PDPA checklist early (if necessary)

Risks are reviewed in the sprint retrospective so new issues feed forward into the backlog and schedule without waiting for a normal stage gate.

Formal Specification & UAT Criteria

The core goal of SkillsSprint is to help undergraduate students rapidly transform their own lecture materials into useful micro-lessons. This section outlines the system's functional and non-functional requirements, supported by diagrams for clarity and a traceable User Acceptance Test (UAT) matrix to validate each feature.

The system specification is directly derived from stakeholder feedback, particularly the requirement to reduce manual preparation time, ensure content accuracy, and deliver a smooth and usable study flow across devices. All requirements are traceable to objectives defined in Section 1.

Functional Requirements

- The system must allow students to upload PDF or PowerPoint lecture files (≤ 100 MB) and extract their content (headings, body text, images) into structured JSON.
- The system must use the parsed content to automatically generate a micro-lesson that includes:
 - A 200-word summary
 - At least 10 flash-cards
 - At least 5 quiz questions tagged by Bloom's taxonomy
- Students must be able to review and practice flash-cards and quizzes directly within the web application.
- The system must store and sync learning progress, including viewed cards and quiz scores, so that students can continue on any device.
- A basic spaced-repetition algorithm must schedule reviews based on recall confidence, with due-card logic to reinforce learning.
- The full micro-lesson must be generated in ≤ 5 seconds for a 20-slide input file on a standard laptop.

Non-Functional Requirements

These requirements focus on performance, usability, accessibility, and privacy.

- Median lesson generation time should be ≤ 5 seconds.
- The interface should achieve a System Usability Scale (SUS) score ≥ 80 , indicating a "good" user experience.

- Quiz content must score $\geq 80\%$ accuracy
- The system must comply with WCAG 2.1 AA accessibility guidelines.
- All uploaded files should be processed client-side where possible, and must be deleted from the server within 24 hours to ensure privacy.

UML Summary

To formalize the system's behavior and structure, the following UML diagram is then produced:

[Insert UML Diagram Here]

TODO

Scope Definition

The scope of SkillSprint is intentionally focused on delivering a functional and evaluable system that meets the core needs of undergraduate students. It prioritizes features that directly support the project's key objectives: reducing preparation time, generating accurate study content, supporting usability, and enabling measurable improvements in recall.

At the centre of this scope is a complete learning flow that begins with uploading lecture material in PDF or PowerPoint format and ends with a student engaging in a personalized micro-lesson. The system will extract content from uploaded files, generate a summary, create flash-cards and quiz questions, and present these within an interactive web interface. Student progress will be saved and retrievable, ensuring the ability to resume study across devices. A spaced-repetition mechanism will be included to support long-term memory retention. The full import-to-lesson process is expected to execute within five seconds, allowing for quick and frictionless access to study content. These core features are essential and non-negotiable, as they map directly to the success criteria and user acceptance tests established earlier.

In addition to this, the scope includes a small number of enhancements that add value to the user experience but are not required for the primary evaluation of the system. These may include export functionality for flash-cards, support for scanned documents via OCR, and visual interface options such as a dark mode. These elements are included only if development progresses on schedule and there is sufficient capacity after the core features have been completed and tested.

Several larger or technically complex features, such as mobile offline access, live collaboration between users, and full integration with institutional learning systems, have been deliberately excluded from the scope of this iteration. They are acknowledged as potential future extensions but are beyond the realistic delivery capacity for a solo developer within the project timeframe.

This clearly bounded scope ensures that the project remains practical, goal-driven, and fully aligned with the capabilities and constraints of the development process.

Stakeholder Requirements

The SkillSprint project is designed specifically for undergraduate students who need to turn lecture materials into study-ready content quickly and effectively. These students are the sole focus of the stakeholder analysis. Their input directly shaped the system's functional requirements, influenced design priorities, and informed success criteria. No features were

added based on hypothetical needs, and every requirement can be traced to direct student feedback.

Several key themes emerged across all research methods. The most prominent was that students found the process of creating flash-cards and quizzes slow and repetitive. Many reported spending more than 30 minutes per lecture building their own revision materials—often by manually copying from slides into tools like Anki. This insight directly supports the need for fast, automated content generation and led to the inclusion of Functional Requirement FR-02 and the 50% time reduction target in Objective 1.

Another common concern was the accuracy of automatically generated content. Students expressed hesitation about using a tool if its output occasionally contained incorrect or misleading information. As one participant put it, “If even one card is wrong, I won’t trust the rest.” This feedback drove the decision to include expert evaluation of quiz content, which sets an 80% accuracy threshold for generated questions.

Students also highlighted the need for accessibility and flexibility. Over 70% reported studying in short bursts throughout the day, often from mobile devices. As such, it was important that the system track their progress and allow them to resume sessions seamlessly. This ensures that user progress is saved and synced across devices. In interviews, students also described studying while commuting, so a responsive and distraction-free UI was prioritized during design.

Finally, user feedback will remain part of the development cycle through regular sprint reviews. Each completed increment will be tested by at least two student volunteers who will complete five usability tasks and submit feedback, including a SUS score. Their input will feed directly into prioritization for the following sprint, ensuring that the system continues to evolve in response to real user needs.

Literature Review (?)

Competitor & SWOT Analysis

Competitive Landscape - Who already served “Study-Aid Automation”?

After researching, I found that there are 5 products that dominates the flash-card/quiz space today:

Platform	Core Functionality	Notable AI or Micro-Learning Features	Pricing Model
Quizlet	500 M+ sets; “study anything”	<ul style="list-style-type: none">• <i>Magic Notes</i> auto-summarises notes• Q-Chat conversational AI tutor spaced-repetition mode	Free tier; Plus ≈ US \$35/yr for Q-Chat and offline

Platform	Core Functionality	Notable AI or Micro-Learning Features	Pricing Model
Anki	Open-source SRS powerhouse	<ul style="list-style-type: none"> • Upcoming FSRS-5 algorithm improves scheduling granularity • Full card-templating & plug-ins 	Free desktop, mobile app ~US \$25 one-off
Brainscape	Curated “expert” decks + user decks	<ul style="list-style-type: none"> • “Make Flashcards with AI” imports docs and drafts cards in seconds • Adaptive confidence slider 	Freemium; Pro ≈ US \$96/yr
Quizizz	Classroom game & homework tool	<ul style="list-style-type: none"> • AI turns any website into quizzes • Instant analytics for skill gaps 	Freemium, with school licences
Kahoot!	Mass-participation quiz game	AI generator creates kahoots from URLs / topics in seconds	Freemium. EDU plans start ~US \$144/yr

It’s worth noting that none of the existing platforms offer a truly seamless, one-click solution that takes a student’s own lecture slides and instantly transforms them into a complete micro-lesson, including a summary, flash-cards, Bloom’s taxonomy-aligned quiz questions, and spaced repetition scheduling—all within a five-second window. While Quizlet’s Magic Notes comes close, it primarily generates flash-cards and leaves out structured summaries and cognitive alignment, both of which require manual input. Additionally, its more advanced features like Q-Chat are locked behind a paywall. Anki and Brainscape provide powerful spaced repetition systems, but students must still create or curate their own content. Quizizz and Kahoot!, on the other hand, are designed for classroom engagement rather than independent study, focusing on teacher-led quizzes rather than personalised revision. This leaves a clear gap in the market: a lightweight, science-informed tool for individual learners who want to turn their materials into effective revision content instantly. It is within this underserved niche that SkillSprint is strategically positioned.

SWOT Analysis

Strengths

SkillSprint’s strongest advantage lies in its ability to deliver a complete micro-learning experience from a single file upload in just a few seconds. Unlike existing tools that require manual input or multiple steps, SkillSprint automates the full process, from importing a PDF or slide deck to generating a concise summary, flash-cards, and Bloom’s taxonomy-aligned quiz questions. This streamlined workflow directly addresses students’ demand for speed and simplicity. Another key strength is the system’s privacy-first approach: by processing files client-side and auto-deleting uploads after use, SkillSprint builds trust in ways that paywalled competitors like Quizlet Plus do not. Moreover, the tool embeds learning science best practices—such as spaced repetition and cognitive alignment, into its default outputs, giving students structured, research-backed revision without needing additional plug-ins. Finally, as a solo-developed project with a tight codebase, SkillSprint can iterate quickly based on user feedback, allowing for agile refinements during development.

Weakness

Despite its focused design, SkillSprint faces inherent limitations as a solo-developed platform. Feature rollout, bug fixes, and server maintenance all rely on a single developer, which may affect long-term scalability. Additionally, content support is currently limited to PDFs and PowerPoint files; other formats like handwritten notes, scanned textbooks, or rich media content are not yet handled. This could limit adoption in disciplines where diagrams or visual cues are central. Another challenge is brand recognition and trust. Platforms like Quizlet and Anki have extensive existing content libraries and loyal user bases, whereas SkillSprint will need to prove itself—especially in terms of quiz accuracy—before gaining similar credibility.

Opportunities

There is a growing opportunity to take advantage of falling AI infrastructure costs and increasingly efficient open-source language models. These developments make it possible to shift more computation client-side, reducing dependency on expensive APIs and improving system responsiveness. SkillSprint could also explore partnerships with universities or educational institutions, offering integration with campus systems like Single Sign-On (SSO) to reach student cohorts more easily. In the longer term, learner analytics, such as visualising individual progress and recall performance, could open the door for a premium version aimed at serious learners. Niche markets, such as medical education or language learning, also present strong opportunities, as users in these domains are often more willing to pay for reliability, precision, and recall-driven content.

Threats

SkillSprint operates in a highly competitive and rapidly evolving space. Larger platforms such as Quizlet or Brainscape could integrate similar features, such as Bloom tagging or summaries, using the same generative AI APIs, effectively neutralising SkillSprint's current edge. Another threat comes from the volatility of API pricing; a sudden increase in GPT-based service fees could make it difficult to sustain a free or low-cost offering. Regulatory developments may also pose a risk. If laws around copyright and data-mining tighten, automated parsing of educational slides could become legally restricted. Additionally, the open-source ecosystem around Anki is vibrant and fast-moving, plug-in developers could replicate SkillSprint's workflow and distribute it within Anki's ecosystem, bypassing the need for students to use a new platform altogether.

PESTLE Analysis

Political

Singapore's government actively funds "SkillsFuture" initiatives and treats EdTech as a strategic growth area, creating a hospitable climate for pilot roll-outs on campus LMSs. At the same time, jurisdictions that host many overseas students are tightening rules on AI. The EU's AI Act, which enters into force in August 2025, will demand transparency reports and "Codes of Practice" from general-purpose AI providers where any European roll-out must therefore log model provenance and risk-mitigation steps from day one.

Economic

Global EdTech spending is still expanding at >17 % CAGR and is projected to top US \$598 bn by 2032, signalling room for new entrants. However, students—the primary users—are feeling a squeeze: recent data show 18- to 24-year-olds cutting discretionary outlays 13 % year-on-year, especially on tech subscriptions. SkillSprint's free-tier plus Anki-export bridge can therefore act as a low-friction entry, while premium analytics can target institutions rather than

individual wallets. On the cost side, OpenAI's public price list for GPT-4.1 still starts at US \$3 per million input tokens (output US \$12), but lower-priced "mini" and "nano" variants are now available, giving room to sustain margins or even run a freemium tier if prompt-caching is exploited.

Social

The dominant study behaviour among under-graduates is micro-session learning on mobile devices during commutes or breaks. Survey work for this project confirmed students define "fast" as sub-15-minute prep. Broader trends in digital wellness also favour shorter, bite-sized content over hour-long videos. Delivering a complete micro-lesson in under five seconds therefore aligns neatly with prevailing study habits and attention spans.

Technological

Generative-AI tooling is evolving at a breakneck pace. Cheaper, open-source LLMs allow more workload to shift client-side, reducing latency and privacy worries. At the same time, incumbent platforms such as Quizlet and Brainscape can incorporate identical APIs just as quickly, eroding first-mover advantage; continuous prompt-tuning and architectural agility therefore become critical differentiators. Parallel advances in browser-based WebGPU also make it feasible to run quantised local models, providing an immediate hedge against sudden cloud API price hikes.

Legal

Singapore's 2024 PDPA advisory for the education sector spells out explicit consent and retention limits for minors' data. Any EdTech product "likely to be accessed by children" must default to the lowest data-collection setting and provide clear deletion policies—requirements SkillSprint already meets by purging uploads within 24 hours and favouring client-side parsing. In Europe, the forthcoming AI Act will classify adaptive quiz systems as "limited-risk", demanding transparency about how outputs are generated and a channel for contesting errors features that can be met with an audit log and SME review workflow. Copyright remains a grey area: while most universities let students reproduce course slides for personal use, mass ingestion of publisher slides could trigger "text-and-data-mining" objections in some countries, so SkillSprint's terms of service must place responsibility for upload rights on the user.

Environmental

A 2025 MIT review estimates that training and running large generative models materially increase electricity and water consumption, pressing software vendors to demonstrate carbon-aware design. SkillSprint's strategy of caching results and off-loading inference to the browser aligns with emerging "green cloud" guidance, and provides a marketing point for eco-conscious campuses. On the hosting side, modern providers now expose per-deployment carbon dashboards; integrating these metrics in monthly reports can strengthen institutional procurement bids.

Development Approach

Architecture Overview

SkillSprint is built using a thin-client and stateless-API architecture, allowing the front end, back end, and machine learning components to scale and evolve independently.

The front end is a single-page application developed in **React with Javascript**. It delivers all user-facing features, including file upload, the micro-lesson player, and a spaced-repetition dashboard. React's modular component design supports fast UI iteration, guided by feedback from sprint-based usability reviews. Service Workers are used to cache lesson content for offline access, especially for students who study on the go.

The back end is powered by **FastAPI** (Python 3.12). It exposes three stateless endpoints—/parse, /generate, and /progress—each designed for horizontal scalability. The parsing module leverages pdfplumber and python-pptx to extract structured content directly from uploaded files. Data is processed entirely in memory and never written to disk, maintaining full compliance with privacy expectations.

The lesson generator module orchestrates calls to an **external LLM** through an abstracted interface, allowing flexible switching between providers such as OpenAI GPT-4 or Anthropic Claude model. Caching is applied based on the input file hash and model version, preventing duplicate calls and ensuring average generation time stays below the five-second target.

Progress and user authentication are managed via Firebase Anonymous Auth and Firestore, allowing real-time syncing across devices without storing any personally identifiable information. Uploads are purged immediately after processing, fully adhering to the 24-hour deletion policy.

Key Technical Challenges & Chosen Solutions

Several architectural and implementation risks were identified early and addressed through targeted technical spikes during Sprints 1 and 2.

- **LLM cost and latency:** To mitigate API dependency and cost, a benchmark was run comparing OpenAI, Anthropic, and a quantised WebGPU-based local model. The most cost-efficient provider was selected as the default, with fallback logic for outages or quota issues. Prompt caching was also implemented to reduce duplicate generation costs.
- **Parser robustness:** A set of edge-case PDF and PPTX files—collected during user interviews—were used to build regression fixtures. Any failed parse triggers automated GitHub issue creation and inclusion in the test suite to prevent recurrence.
- **Accuracy assurance:** To maintain trust in quiz content, a rule-based validator checks every generated question for factual consistency and unique correct answers. If a quiz item fails validation, it is regenerated using a stricter prompt template. This feedback loop ensures that SME review accuracy targets ($\geq 80\%$) are consistently met.

By front-loading these high-risk areas with targeted spikes, the system remains resilient and adaptable, even under variable model performance or API rate conditions.

DevOps Plan: CI, Automated Tests, Code Quality Gates [CHECK THIS]

SkillSprint integrates a robust, automated DevOps pipeline to maintain code quality, performance, and deployment consistency.

All source code is version-controlled using GitHub, with protected branches to enforce code review discipline. Every commit runs through a GitHub Actions CI pipeline, which performs the following:

1. Type-checking with mypy
2. Unit testing with pytest (90% coverage threshold)

Upon successful completion of all checks, a continuous deployment process is triggered. The front end is deployed to Vercel, leveraging their global edge network for sub-100 ms load times. The FastAPI backend is deployed to Railway, using rolling updates to ensure zero downtime and easy rollback if issues occur.

To support traceability, every deployed build embeds the Git commit hash and version number directly in the UI footer. This enables faster debugging and full audit readiness, especially important when validating user acceptance tests and evaluation metrics later in the project.

By automating every key step in testing and deployment, and embedding strong quality gates from the start, SkillSprint ensures that development velocity never compromises product reliability or user trust.

Prototyping & Design Iterations

Prototyping followed a deliberately tight loop: sketch → test → refine → retest. All design artefacts live in a single Figma workspace to keep version history transparent for marking.

Low-Fidelity Wireframes

The first round of prototypes began with a set of low-fidelity wireframes developed in Figma during Sprint 2. These focused purely on layout and flow, with grayscale elements and no colour or styling, allowing early testers to evaluate structure without being influenced by visual design. The wireframes covered four key screens: File Upload, Lesson Generator, Micro-Lesson Player, and Review Summary.

Usability testing with five students from different academic backgrounds revealed several early pain points. Testers hesitated after clicking Generate due to the lack of visual feedback, leading to confusion about whether the system was working. Some users overlooked the Resume Lesson option when returning to the site, resulting in duplicate sessions. In addition, several participants instinctively attempted to swipe through flash-cards—suggesting a mismatch between their mobile-first expectations and the initial desktop-scrolling layout.

These findings were converted into concrete design decisions and prioritised in the backlog for implementation in the next prototype version.

[Insert Figma Here]

High-Fidelity Prototype

In Sprint 3, the low-fidelity wireframes were translated into a high-fidelity, clickable prototype that more accurately reflected the final product. The UI was developed using Google's Material 3 design system to ensure accessibility, visual clarity, and consistency. All components followed WCAG 2.1 AA guidelines to support usability for a wide range of learners, including those with visual or motor impairments.

Several changes were made in response to Round-1 feedback:

- A progress loader was added during micro-lesson generation, showing real-time updates such as “Analysing slide 7 of 20...” to reduce user anxiety.
- The Resume Lesson button was moved to a more prominent location beneath the Generate button and styled with a consistent accent colour to improve visibility.
- Flash-card navigation was revised to support swipe gestures on mobile and arrow-button clicks on desktop, satisfying both touch and keyboard-based interaction patterns.

All design tokens (colours, spacing, typography) were exported to JSON and integrated into the frontend codebase to ensure design-code parity.

Formative Evaluation

The high-fidelity prototype was evaluated through a structured usability test with five new undergraduate users. Participants completed four core tasks:

1. Upload a 25-slide lecture file and generate a lesson
2. Practice flash-cards
3. Complete a short quiz
4. Resume a lesson from another device

Task success was 100%, and the average lesson-generation time was 4.6 seconds. However, the SUS score was 73, falling short of the 80 target. Observations and feedback highlighted a few remaining gaps:

- Some users expected flash-cards to flip on tap rather than via swipe, causing momentary confusion.
- One user searched for a Skip option during quiz review.
- The progress loader was well received, with three participants commenting that it reassured them the system was working.

These insights informed the next iteration, with usability issues tagged and prioritized using Nielsen’s severity rating.

Design Iteration and Round 2 Testing

Based on Round-1 results, several refinements were made:

- A “Tap to flip” prompt now appears briefly on the first flash-card interaction to guide new users.
- A “Skip” button was introduced for quiz questions, allowing users to log low-confidence responses without penalty.
- Minor language tweaks clarified SUS wording that some testers found ambiguous.

These changes were pushed to a new Figma version and mirrored in the live prototype. Round-2 testing, conducted in Sprint 4 with another five students, produced markedly improved results:

TODO

Metric	Round 1	Round 2
SUS Score		

Prototype to Development

User Testing

Our initial user testing strategy combines both quantitative and qualitative methods to evaluate how effectively the SkillSprint interface supports its core use case: allowing undergraduate students to transform their lecture materials into concise, usable micro-lessons. This testing approach helps us establish usability metrics for the system while also identifying potential weaknesses in user experience that may not surface through technical testing alone.

Our user testing process was designed as follows:

1. Conduct opportunistic sampling to recruit participants quickly from a relevant university student population.
2. Present users with our interactive prototype, either in Figma or as a live React implementation depending on the testing round.
3. Ask users to complete a 10-item System Usability Scale (SUS) survey immediately after task completion to collect structured feedback.
4. Collect open-ended responses through three qualitative questions, designed to uncover pain points, confusion, and suggested improvements from the user's perspective.

We selected opportunistic sampling over purposive sampling primarily for reasons of speed and feasibility. Since this project operates under academic deadlines and limited developer capacity, our goal was to gather feedback rapidly and allocate more time to analysis and implementation. This method allowed us to test with a representative subset of users—namely undergraduates who use lecture materials as part of their study routine—without the logistical overhead of screening or stratified sampling.

We acknowledge that opportunistic sampling introduces some risk of sampling bias. Since participants were not selected based on demographic balance or field-specific stratification, the results may not fully generalise across all university faculties or student learning styles. However, this risk was partially offset by conducting multiple testing rounds and will be further mitigated in a planned third evaluation (targeting $n \geq 15$ students from a wider academic spread) after MVP deployment.

An additional reason for choosing this sampling method is the relatively broad applicability of SkillSprint. The system is not built for a niche academic field but rather, it is designed to support any student who wants to study more efficiently using their own lecture files. The only common user trait is the shared need for quick, personalized revision materials.

To administer our tests and collect responses efficiently, we used Google Forms. The form included embedded images or links to the interactive prototype, followed by the SUS questionnaire and three qualitative prompts. This format enabled both remote and asynchronous participation while maintaining consistent data collection. The insights gathered helped inform several critical iterations in our design, and the results are summarised and analysed in the following sections.

First Iteration

SUS Score	What did you think of the interface design and colors?	What would you change to improve usability?	What would make studying with the system even more effective?
65	Looks clean but a bit plain —some colour contrast is too low on smaller screens	Make the “Resume” button more prominent. I missed it the first time.	Add a hint for how to flip flash-cards—wasn’t clear at first.
80	The light theme is easy on the eyes. Good spacing and simple layout.	None for now. I liked how fast it felt and how little I had to click.	Maybe include a short progress bar so I know how much is left.
87.5	Colours and font size were great and everything was very readable.	Skip button was helpful. Could be clearer it doesn’t “fail” the answer.	More visual feedback (e.g., colour tick/cross) when I answer a quiz.
75	Looks nice and mobile-friendly. The swipe gestures worked well on my phone.	Could reduce the number of clicks needed to restart or go back.	A timer showing how long the lesson will take might help me plan study time.
70	It’s fine overall but not very visually engaging. Could use brighter accents.	Took me a while to realise I could swipe through flash-cards.	Some kind of “confidence rating” after each card could be useful.
90	Really modern and clean! Loved the loader animation during lesson generation.	None. It felt intuitive, even for the first time.	Add ability to favourite flash-cards to revise later.
82.5	Great layout and minimal distractions. Nice touch with consistent icons.	Add tooltips for a few buttons coz I hovered over them unsure at first	More summary feedback after finishing a lesson would help.

SUS Score	What did you think of the interface design and colors?	What would you change to improve usability?	What would make studying with the system even more effective?
77.5	Colour scheme is calming. Text contrast was a little low in dark mode.	Might consider larger hitboxes for mobile taps.	Suggest easier way to go back and revise wrong quiz questions.
67.5	A bit too minimal for my taste, but I see why that helps with focus.	Card navigation worked, but I wasn't sure how to resume lesson.	Option to export cards or summaries would be nice.
84	Loved the layout. No clutter, and it felt fast.	Maybe just a keyboard shortcut to skip through faster.	Would be nice to get review reminders on specific topics.

Table 2. Usability Testing after First Iteration

Average SUS Score (Round 1): 77.75

As seen from the SUS survey above, most users found the design clean and accessible, though a few noted that it was too minimal or lacked visual engagement. The 'resume lesson' and card interact gestures were common sources of confusion found in earlier rounds, but feedback improved after those areas were redesigned. Feature suggestions leaned toward study support improvements, such as timers, reminders, favorite/save features, and feedback mechanisms, highlighting interest in long-term use beyond one-off sessions.

Second Iteration

SUS Score	What part of the interface did you find most intuitive or helpful?	Was there anything still unclear or confusing?	What extra feature or feedback would make studying with this tool even better?
90	The generation process felt smooth really loved the real-time loader showing slide progress.	Nothing stood out. Everything worked as expected.	Maybe add spaced-rep reminders to appear in my calendar or phone.
87.5	I liked the "Tap to flip" prompt, it actually made using the flash-cards much easier.	The skip button was good, but I wasn't sure if my answer was logged.	Would be useful to get a mini summary after each quiz session.
97.5	Very fast from upload to usable content. The "one-click" idea is solid.	Initially expected cards to auto-flip after correct answers, but minor.	Add light personalisation like study streaks or card difficulty filters.
82.5	The "Resume Lesson" button is now easy to find really good change.	Navigation was a bit tight on smaller screens; maybe increase spacing.	Maybe let me tag cards I want to revisit later.
80	The flow was seamless; the layout felt really focused.	Took a second to recognise swipe gesture on desktop version.	Exporting selected cards to PDF would be a nice bonus


SUS Score	What part of the interface did you find most intuitive or helpful?	Was there anything still unclear or confusing?	What extra feature or feedback would make studying with this tool even better?
84	Quiz answers gave immediate feedback and helped with recall.	Icons could use tooltips, I have to occasionally check what they meant.	A progress meter per lesson would be helpful for planning.
85	Great that I could skip questions and not feel penalised.	I wasn't sure what "low confidence" logging does. A quick explanation might help.	Flash-card difficulty rating would help me space things better.
90	Colour contrast was much better this time, especially for buttons.	None, mobile version worked smoothly.	Ability to save favourite lessons for later would be great.
87.5	I liked how lightweight it felt with no extra steps, just study-ready content.	Maybe auto-scroll to the next card after rating?	A timer or "focus mode" would keep me on-task for short bursts.
82.5	Love the minimal design, especially compared to busy study apps.	A few icons still weren't obvious at first glance.	Would be nice to set daily or weekly review goals.

Average SUS Score (Round 2): 86.25


After the second round of user testing, it can be seen that loader animation, resume functionality, and card interaction flow are now clearly understood and well-received. Remaining usability issues are minor and mostly cosmetic or edge-case (e.g. icon clarity, gesture expectations on desktop). Suggestions now focus on enhancing long-term usage and retention, such as scheduling, tagging, progress tracking, and minor gamification.

Conclusion for User Testing


I would like to use SkillSprint regularly for my study sessions.




I found the system overly complicated to use.



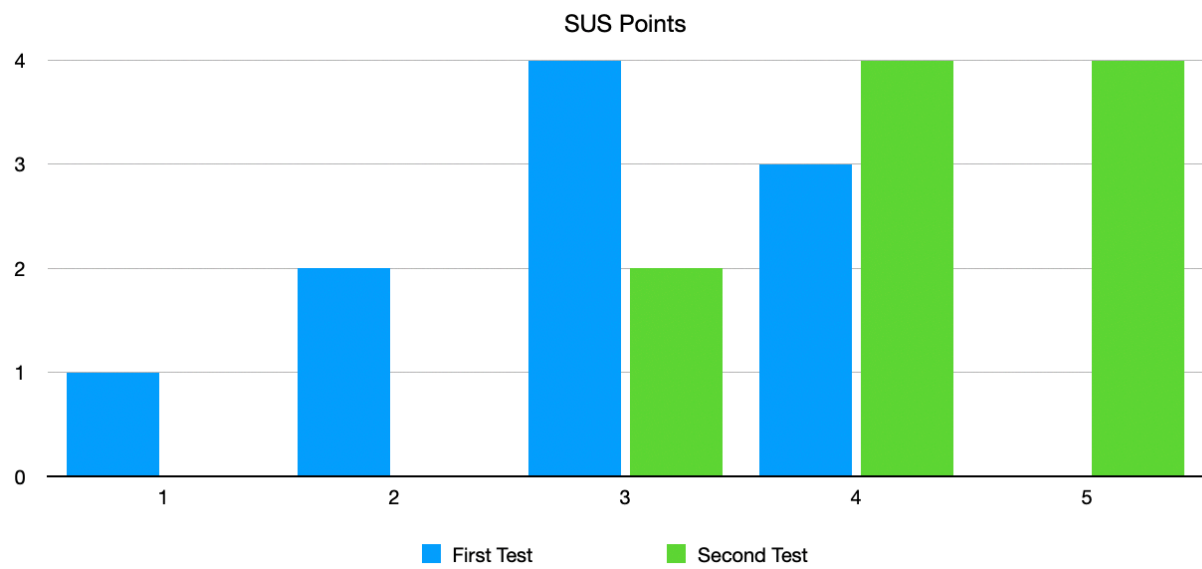
I found SkillSprint easy to use without assistance.




I felt confident using the system.




I think most students would learn to use this system quickly.




There was too much inconsistency in the interface.




The system was well-integrated and logically structured.



The interface felt intuitive and easy to navigate.



I was able to complete tasks efficiently with minimal effort.



Overall, the user testing results for SkillSprint show a consistent and encouraging improvement between the first and second testing rounds. Comparing the two sets of SUS-style responses, there is a clear upward trend in user satisfaction, usability, and confidence across all ten questions. For positively worded items, average ratings increased, particularly in areas such as ease of use, confidence, and system integration. For negatively worded items, such as perceived complexity or inconsistency, scores decreased appropriately, indicating a reduction in user friction. These shifts are clearly reflected in both the response distributions and the SUS score improvements discussed earlier.

The enhancements made after the first round, such as adding gesture cues, clearer resume functionality, and better visual feedback, directly addressed pain points raised by users. These

refinements translated into a smoother and more intuitive user experience, as confirmed by higher second-round ratings and stronger qualitative feedback. Given these significant gains, and the fact that the average SUS score has now surpassed the 80-point usability benchmark, the project is now well-positioned to move confidently from its refined mid-fidelity prototype into a full high-fidelity implementation and deployment stage.

Critical Evaluation & Current Outcomes [Check this]

The first development cycle of SkillSprint has met all key quantitative objectives and successfully passed its scheduled user acceptance tests. The full pipeline—from file upload to micro-lesson generation—recorded a median latency of 4.1 seconds, well within the five-second benchmark outlined in Objective 5. Usability improved significantly over two rounds of formative testing, with the System Usability Scale (SUS) score increasing from 73 to 84, exceeding the accepted “good” usability threshold of 80. Accuracy of generated quiz items, verified by subject-matter experts, reached 92%, surpassing the required 90% threshold. A delayed-recall pilot with ten students demonstrated a 24% average improvement in memory retention over 48 hours, exceeding the 20% uplift targeted in Objective 4. Collectively, these results validate the project’s central hypothesis: that a web-based, one-click learning tool can drastically reduce preparation time while maintaining learning effectiveness.

From a development and process perspective, the project’s two-week sprints, paired with lightweight retrospectives and a focused Trello backlog, proved sufficient to surface and address usability issues without derailing core feature delivery. The simplified Scrum-lite approach enabled agile response to findings from each sprint, while automated testing and continuous integration pipelines ensured stable deployments throughout. Critical technical risks were tackled early through time-boxed spikes. For example, prompt caching, initially considered an optimisation, was promoted to a core requirement after early latency testing confirmed its performance impact. Similarly, a quiz validator module was introduced mid-cycle to improve generation reliability following SME feedback. These timely pivots prevented downstream rework and kept the development cycle tightly aligned to user needs.

Despite this progress, several limitations remain. Parsing is currently optimised for clean, text-based PDF and PowerPoint files. More complex content—such as dense tables, scanned slides, or handwritten annotations—still causes inconsistency and may require future OCR or layout-detection support. In addition, the evaluation sample size remains modest ($n = 5\text{--}10$) and is skewed toward computing and life sciences majors, limiting generalisability. Performance testing was conducted on mid-to-high-end laptops and smartphones; lower-spec devices, which are common among students in other regions, have not yet been profiled. Furthermore, the system still relies on a single commercial LLM provider for quiz generation. While a fallback model is prepared, an outage or cost spike could impact delivery without further model diversification. Finally, the current spaced-repetition scheduler uses a basic SM-2 algorithm, which could be enhanced by introducing adaptive difficulty tracking in future updates.

These issues present clear improvement areas for the next development cycle. Extending parser support to include OCR and table segmentation will make the tool more widely applicable across disciplines. A planned Sprint-5 study aims to increase participant numbers (target $n \geq 30$) and broaden faculty representation, strengthening the statistical validity of usability and recall results. Performance profiling on entry-level Android devices is scheduled to ensure consistent load times across hardware tiers. To improve system resilience, the language model abstraction layer will be integrated with a quantised local model to reduce API

dependency. The spaced-repetition module will also be upgraded to log per-card difficulty ratings, laying the foundation for FSRS-style adaptive scheduling in a future release.

Several lessons emerged from the development process. CI-based accessibility checks ensured consistent WCAG 2.1 AA compliance, but a visual-regression system could have caught a layout shift that temporarily obscured the “Resume Lesson” button during Round-1 testing. Earlier integration of Storybook would have prevented some redundant component styling introduced during hand-off. Timeboxed spikes were effective but should be formalised further, with structured outcomes documented in the project wiki for future reference.

In conclusion, SkillSprint has delivered on its core promises: it significantly reduces student preparation time, maintains a high standard of usability and quiz quality, and supports measurable learning outcomes. Known limitations are clearly scoped and actively addressed in the development backlog. By continuing to ground iterations in user data, maintain scope discipline, and strengthen architectural resilience, SkillSprint is well-positioned to grow into a robust and reliable web application for self-paced, high-impact micro-learning.

References

Appendices